

## Lecture 11.3: Efficient LLM Deployment

Quantization, QLoRA, and Running Models at Home

Heman Shakeri

# The Democratization Challenge

From billion-parameter models to your laptop

## Topics:

- ① The memory bottleneck: Why 7B models need 66+ GB for training
- ② Quantization fundamentals: Representing weights with 4 bits instead of 16
- ③ QLoRA: Fine-tuning 7B models on consumer GPUs (8–12 GB)
- ④ Hands-on: Fine-tune and deploy models on your machine
- ⑤ Trade-offs: Speed, quality, and hardware constraints

## Practical Outcome

By the end, you will fine-tune and serve models on consumer hardware!

# The Billion-Parameter Reality

Modern language models are remarkably capable but extraordinarily large:

## Open models:

- **Gemma 3:** 270M, 2B, 9B, 27B (Apache-2.0)
- **LLaMA 3.1:** 8B, 70B, 405B (Apache-2.0)
- **Mistral:** 7B (Apache-2.0)
- **Qwen 2.5:** 7B, 14B, 72B (Apache-2.0)

## Proprietary models:

- **GPT-4:** Rumored 1.7 trillion parameters
- **Claude 3:** Multiple variants (sizes undisclosed)

**Question:** Can you run these on your laptop? Let us do the math.

## Memory Requirements: Just Storing Weights

Each parameter is a number. Standard storage formats:

Format	Bits/param	Bytes/param	7B model
FP32 (float32)	32	4	28 GB
FP16 (float16)	16	2	14 GB
BF16 (bfloat16)	16	2	14 GB
INT8	8	1	7 GB
INT4	4	0.5	3.5 GB

**Example:** 7 billion parameters in FP16:

$$7 \times 10^9 \text{ params} \times 2 \text{ bytes/param} = 14 \text{ GB}$$

This is just the **storage cost**. **Inference** requires *more* memory to hold the **KV-Cache** (the “memory” of the conversation), which can be several GBs itself for long contexts.

**Your hardware:**

- Consumer laptop: 8–16 GB RAM (shared with OS, browser)
- RTX 3060/4060 GPU: 8–12 GB VRAM
- MacBook Pro M3: 16–32 GB unified memory

## What About Smaller Models?

Gemma 3 270M in different precisions:

Format	Inference	Training (w/ Adam)
FP32	1.08 GB	13 GB
FP16	540 MB	6.5 GB
INT8	270 MB	–
INT4	135 MB	–

**Much more manageable!** But what if you want to use larger, more capable models?


**Trade-off:**

- 270M: Fits easily, limited capability
- 7B: Better capability, needs optimization
- 70B+: State-of-art capability, requires advanced techniques

**Today's goal:** Make 7B models as accessible as 270M models used to be!

## Training Makes It Worse

During training, memory requirements explode:



### Why so much?

- **Gradients:** Same size as weights (14 GB for FP16)
- **Optimizer states:** Adam stores first and second moments in FP32 (28 GB!)
- **Activations:** Stored during forward pass for backpropagation (10+ GB)

This requires enterprise GPUs (A100 80GB, H100 80GB) costing \$10,000–\$40,000!

# The Challenge and Solution

## The Challenge

**Goal:** Enable researchers and practitioners to:

- **Run** 7B models for inference on 8 GB GPUs
- **Fine-tune** 7B models on 12 GB GPUs
- **Deploy** models on CPUs and consumer hardware

## The Solution

**Three-pronged approach:**

- ➊ **Quantization:** Fewer bits per parameter
- ➋ **Efficient fine-tuning:** QLoRA (quantization + LoRA)
- ➌ **Optimized serving:** Specialized inference engines

**Result:** Democratize access to powerful LLMs!

## The Core Idea

**Quantization:** Represent weights with fewer bits while preserving performance

**Analogy:** Image compression

- Raw photo: 24 bits per pixel (16 million colors)
- JPEG: Approximate values, much smaller file
- Visual quality: Often indistinguishable


**For neural networks:**

- Original weights: FP16 or FP32 (continuous values)
- Quantized weights: INT8 or INT4 (discrete levels)
- Model quality: Minimal degradation with proper techniques!

**Key insight:** Neural networks are remarkably robust to reduced precision

# How Quantization Works

**Basic quantization** maps floating-point weights to integers



**The quantization formula:**

Given floating-point weight  $w \in \mathbb{R}$ , quantize to integer:

$$w_{\text{quant}} = \text{round} \left( \frac{w - z}{s} \right)$$

where  $s = \text{scale}$ ,  $z = \text{zero-point}$

To reconstruct (dequantize):

$$\hat{w} = w_{\text{quant}} \cdot s + z$$

## Example: 8-bit Quantization

**Task:** Quantize weights in range  $[-0.5, 0.5]$  to INT8 (256 values)

**Steps:**

- 1. Determine range:**  $w_{\min} = -0.5$ ,  $w_{\max} = 0.5$
- 2. Compute scale:**

$$s = \frac{w_{\max} - w_{\min}}{2^8 - 1} = \frac{1.0}{255} \approx 0.00392$$

- 3. Quantize  $w = 0.12$ :**

$$w_{\text{quant}} = \text{round}\left(\frac{0.12}{0.00392}\right) = \text{round}(30.6) = 31$$

- 4. Dequantize:**

$$\hat{w} = 31 \times 0.00392 \approx 0.1215$$

**Error:**  $|0.12 - 0.1215| = 0.0015$  (very small!)

# Quantization Strategies

Method	Description	Bits
PTQ	Post-Training Quantization	4–8
QAT	Quantization-Aware Training	4–8
Weight-Only	Quantize weights, keep activations FP16	3–4
Weight + Act	Quantize both	8

For LLMs, weight-only quantization dominates:


- Quantize weights to INT4 (16 values)
- Keep activations in FP16 during computation
- **4× memory reduction** with minimal quality loss!

Weight-only works because weights ( $W$ ) are static and have a nice, stable distribution. Activations ( $x$ ) are dynamic (change with every input) and have large outliers. Quantizing these outliers (e.g., in  $W_x$ ) is what causes performance degradation.

## NF4: Normal Float 4

Standard INT4 uses uniform spacing:  $\{0, 1, 2, \dots, 15\}$

But neural network weights follow a **normal distribution!**



NF4 (information-theoretically optimal)

More levels where  
there is more data

**NF4** (used in QLoRA): More quantization levels near zero, fewer in tails


This gives **better precision where it matters!**

# Quantization Methods for LLMs

Method	Description	Use
GPTQ	Post-training using Hessian info	3 – 4-bit inference (static quantization)
AWQ	Activation-aware (protects salient weights)	Fast inference
GGUF	Format for llama.cpp (CPU/Metal)	Laptops, no GPU
bitsandbytes	On-the-fly 4/8-bit with NF4	Training (QLoRA) & inference

**Bottom line:** With 4-bit quantization, the **16-bit model size is reduced by  $\approx 4 \times$ :**  
e.g. 70B model:  $\sim 140 \text{ GB} \rightarrow 41 \text{ GB}$ .

**Quality Loss:** 4-bit QLoRA with NF4 is shown to **match 16-bit performance** [Fig2 of QLoRA paper:].




# QLoRA

Recall Memory savings for adapters:

- Original trainable params:  $d \times k$
- LoRA trainable params:  $r(d + k)$  (typically  $r = 8$  or  $r = 16$ )

**The problem:** Even with LoRA, the **65B** model still requires **> 780 GB** in **16-bit** for finetuning!

**QLoRA** = Quantized base model + LoRA adapters



# QLoRA: Visualizing the Memory Savings




Figure 1: Different finetuning methods and their memory requirements. QLoRA improves over LoRA by quantizing the transformer model to 4-bit precision and using paged optimizers to handle memory spikes.

- Full Finetuning updates a huge **16-bit** model and stores a massive **32-bit** Optimizer State.
- LoRA reduces the Optimizer State slightly by adding small **16-bit** Adapters, but the base **16-bit** model remains the memory bottleneck.
- QLoRA fundamentally solves the bottleneck by quantizing the base model to **4-bit**, while retaining the small **16-bit** Adapters and using Paged Optimizers (Paging Flow) to manage the Optimizer State with the **CPU**.
- *Paging* is an old memory management scheme that allows a system to store and retrieve data from secondary storage and the main memory

## QLoRA: Key Innovations

### 1. 4-bit NormalFloat (NF4): Quantize base to 4 bits optimally

- Information-theoretically optimal for normally distributed weights
- Better empirical results than 4-bit Integers and 4-bit Floats

### 2. Double Quantization (DQ): Quantize the quantization constants

- Quantization constants are typically 32-bit floats
- DQ quantizes these constants, saving an average of **0.37** bits per parameter
- Saves  $\sim 3$  GB for a **65B** model

### 3. Paged Optimizers: Handle memory spikes via unified memory

- Uses **NVIDIA** unified memory for automatic page-to-page transfers between **CPU** and **GPU**
- Prevents out-of-memory errors during gradient checkpointing peaks

### 4. LoRA adapters in BF16: Keep trainable parameters in high precision

- The frozen base is 4-bit; the trainable adapters are **16-bit BrainFloat (BF16)**
- Training with high precision is critical for stable gradient updates

## QLoRA for Gemma 3 270M

Can we use QLoRA for smaller models?

Yes, but it's overkill!

Method	Memory	When to Use
Full FT (FP16)	6.5 GB	Plenty of VRAM
LoRA (FP16)	4 GB	Save some memory
QLoRA (4-bit)	2.5 GB	Extreme constraints

For Gemma 3 270M:

- Regular LoRA already very efficient
- QLoRA adds complexity without much benefit
- Only use if you have <4 GB VRAM

Sweet spot for QLoRA: 7B–70B models on consumer GPUs

For larger models (70B+):

- 70B in 4-bit: ~35 GB (still needs high-end GPU)
- May need model parallelism across multiple GPUs

## Beyond Training: Efficient Inference

You've fine-tuned your model. Now what?

**Deployment requirements:**

- Low latency (time to first token <100ms)
- High throughput (many requests/second)
- Memory efficiency (serve multiple users)
- Hardware flexibility (GPU, CPU, edge devices)

**Challenge:** Standard HuggingFace Transformers is not optimized for production

**Solution:** Specialized inference engines

# Inference Engine Landscape

Engine	Specialty	Hardware
vLLM	High-throughput GPU serving	NVIDIA GPU
llama.cpp	CPU/edge deployment	CPU, Metal, CUDA
MLX-LM	Apple Silicon optimization	M1/M2/M3
TGI	Production-ready API server	NVIDIA GPU
TensorRT-LLM	Maximum GPU performance	NVIDIA GPU

## Key innovation across all engines:

- Quantization (INT4/INT8)
- KV-cache optimization
- Batching and scheduling
- Kernel fusion and optimization

# The Democratization Vision

## The Future is Accessible

**Three years ago:** Running LLMs required million-dollar infrastructure

**Today:** With the techniques you learned:

- Run 7B models on \$300 consumer GPUs
- Fine-tune on laptops with 16 GB RAM
- Deploy on edge devices and phones

**Tomorrow:** Even more efficient methods will emerge!

- Better quantization (sub-4-bit without quality loss)
- Sparse models (Mixture-of-Experts at edge)
- Hardware-software co-design

**Your opportunity:** Use these tools to build LLM applications that were impossible just a few years ago!

## Real-World Impact

The techniques you learned today are **actively used in production**:

**Startups:** Deploy LLM services without massive GPU clusters

**Research labs:** Enable more scientists to experiment with large models

**Edge AI:** Run models on robots, drones, IoT devices

**Privacy applications:** Keep sensitive data local with on-device inference

**Developing regions:** Access AI without expensive cloud services

*“The best model is not the largest model—  
it is the model you can actually deploy and use.”*

# Further Reading

## Core papers:

- Dettmers et al. (2023). "QLoRA: Efficient Finetuning of Quantized LLMs"
- Frantar et al. (2022). "GPTQ: Accurate Post-Training Quantization for GPT"
- Lin et al. (2023). "AWQ: Activation-aware Weight Quantization"
- Kwon et al. (2023). "Efficient Memory Management for LLM Serving" (vLLM)
- Dettmers et al. (2022). "LLM.int8(): 8-bit Matrix Multiplication"

## Tools and frameworks:

- vLLM: [docs.vllm.ai](https://docs.vllm.ai)
- llama.cpp: [github.com/ggerganov/llama.cpp](https://github.com/ggerganov/llama.cpp)
- AutoAWQ: [github.com/casper-hansen/AutoAWQ](https://github.com/casper-hansen/AutoAWQ)
- MLX: [ml-explore.github.io/mlx](https://ml-explore.github.io/mlx)
- bitsandbytes: [github.com/TimDettmers/bitsandbytes](https://github.com/TimDettmers/bitsandbytes)
- PEFT library: [huggingface.co/docs/peft](https://huggingface.co/docs/peft)

## Emerging techniques:

- GPTQ-R: Refined GPTQ with better Hessian approximation
- QuIP: Quantization with Incoherence Processing (2-bit!)
- SpQR: Sparse Quantized Representation
- SmoothQuant: Migrate difficulty from activations to weights