

Lecture 12.2: GenAI: Diffusion Models

From Discrimination to Creation

Heman Shakeri

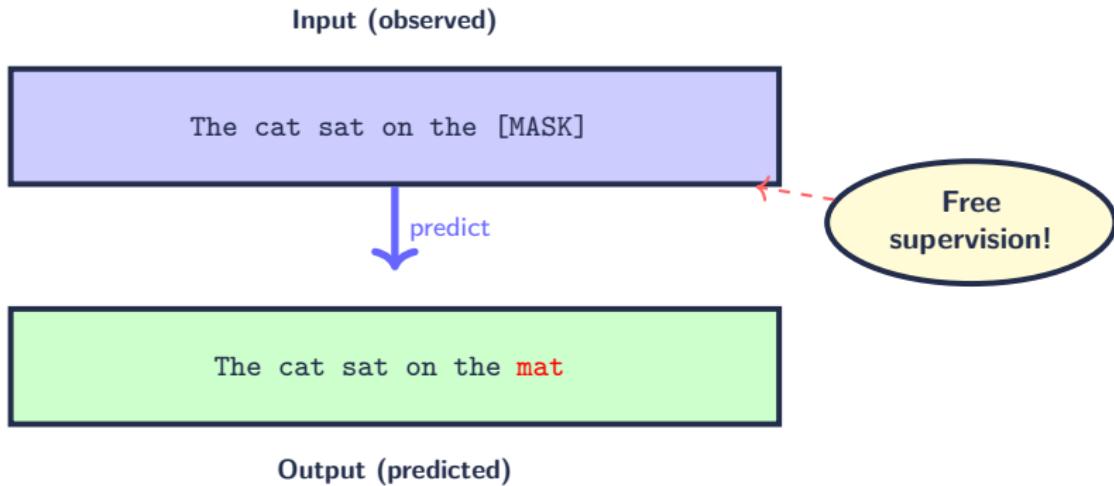
Self Supervised Learning

Most intelligence comes from unsupervised observation

- Babies learn how the world works largely by **observation**
 - Object permanence, gravity, intuitive physics
 - No explicit labels needed
- Humans learn to drive with ~20 hours of practice
 - Leverage vast background knowledge from observation
 - Not millions of labeled examples
- **Common sense:** Generalized knowledge about the world
 - Taken for granted in humans
 - The “dark matter” of AI (LeCun & Misra, 2021)

Self-Supervised Learning: Recall the Core Idea

Learn to predict hidden parts from visible parts



Key insight: The data itself provides the training signal

- No manual labeling required
- Can scale to billions of examples

Self-Supervised Learning for Language

Two dominant strategies from Lectures 9–11:

Strategy	How it works	Examples
Masked Language Modeling	Mask 15% of tokens, predict them	BERT, RoBERTa
Autoregressive	Predict next token given all previous	GPT, LLaMA

Why this works for text:

- **Discrete tokens:** Can enumerate all possibilities
- **Manageable length:** Hundreds to thousands of tokens
- **Natural ordering:** Left-to-right for autoregressive
- **Compactness:** Can represent probability over entire vocabulary

We observed that these models learn rich semantic representations without labels!

Auto-regressive: One Step at a Time

Key Insight: We can re-frame the autoregressive for vision: Instead of regressing a blurry “average” pixel, we can predict a **probability distribution** over *discrete* pixel values (e.g., 0-255).

Process: (Factorization of the joint distribution)

- ① Start with image missing all pixels
- ② Predict first pixel **class**: $p(x_1)$ (Softmax over 256 values)
- ③ Predict second given first: $p(x_2 | x_1)$
- ④ Continue: $p(x_t | x_{<t})$ for all t

Factorization:

$$p(\mathbf{x}) = p(x_1) \cdot p(x_2 | x_1) \cdot p(x_3 | x_1, x_2) \cdots p(x_T | x_{<T})$$

Add diversity: Sample from the discrete distribution $x_t \sim p(x_t | x_{<t})$

Auto-regressive: Successes and Limitations

Modern successes:

- **GPT models** (including ChatGPT): Text generation, one token at a time
- **PixelCNN**: Image generation (Masked Convolution)
 - It uses Softmax over 256 pixel values!
 - Avoids blur by treating pixels as discrete classes, not continuous.
- **WaveNet**: Audio generation (Dilated Convolution)
 - Same trick: Uses Softmax over 256 *quantized* audio levels.

The Problem for Images:

Modality	Length	Auto-regressive?
Text (GPT)	Thousands of tokens	Efficient!
Images (512×512)	262,144 pixels	Too slow!

Challenge: The discrete approach works, but is computationally infeasible.
Can we find a new way to handle continuous pixels **in parallel**?

The Vision Challenge: Why Not Just “Tokenize” Images?

Why can't we just apply text SSL (BERT/GPT) strategies to images?

e.g., Break image into 16x16 patches (like ViT) and treat them as “tokens”?

■ Problem 1: No Finite “Token” Vocabulary

- Text: We have a shared, discrete vocabulary ($\sim 50K$ tokens). We can use **Softmax**.
- Images: A “patch” is a high-dimensional **continuous vector** ($16 \times 16 \times 3 = 768$ dims).
- **We cannot run Softmax over an infinite, continuous space!**

■ Problem 2: The Averaging Problem Returns

- “Okay, so let's predict the continuous patch *vector* using **Regression** (MSE loss).”
- **This fails!** The **target** (the patch) is a set of highly **correlated** pixels.
- If a patch has multiple valid completions (e.g., pointy ear, floppy ear), the MSE loss forces the model to predict the **average vector**.
- **Result:** A blurry, unrealistic patch. **Predicting a correlated target fails!**

■ Problem 3: No Natural Ordering

- Text has a clear 1D (left-to-right) structure for autoregression.
- Images are 2D. A raster scan (row-by-row) of patches is arbitrary and inefficient.

The Key Insight: Change The Prediction Target

The Problem with Masking: The target (the masked patch) is a vector of **correlated** pixels. Predicting it with MSE causes the **averaging problem**.

The Solution (Diffusion): Change the task! Instead of predicting the *patch*, predict the *noise* we added.

Why This Works (Denoising):

- ① **New Target:** The target is now the **Gaussian noise vector** ϵ .
- ② **Independence:** By definition, the noise ϵ is **uncorrelated** across all pixels.
- ③ **Averaging Solved:** We can now use a simple MSE loss ($\|\epsilon - \hat{\epsilon}\|^2$) to predict all 786,432 independent noise values **in parallel!**
- ④ **SSL Signal:** This is a perfect SSL task: we know the noise we added, so we have the ground truth for free.

Predicting independent noise avoids the averaging problem!

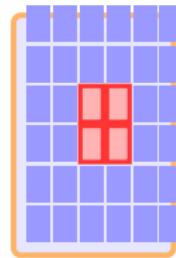
Visual Comparison: Three Strategies

Autoregressive



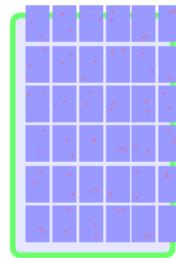
1 pixel at a time
262K steps
Too slow!

Masking



Predict patch
~1K steps
Blurry!

Diffusion



All pixels + noise
50-100 steps
Optimal!

Diffusion: $> 2000\times$ speedup over autoregressive for images!

The Diffusion Process: Overview

Two complementary processes:

1. **Forward (Fixed)**: Gradually add noise over T steps

- Start: Clean image \mathbf{x}_0
- End: Pure Gaussian noise $\mathbf{x}_T \sim \mathcal{N}(0, \mathbf{I})$
- *This is a fixed, known process*

2. **Reverse (Learned)**: Gradually remove noise

- Start: Pure noise \mathbf{x}_T
- End: Clean image \mathbf{x}_0
- *This is what we learn with a neural network*

If we know how to reverse the noising, we can generate!

Forward Process: Adding Noise

Iterative formulation:

$$\mathbf{x}_t = \sqrt{1 - \beta_t} \cdot \mathbf{x}_{t-1} + \sqrt{\beta_t} \cdot \boldsymbol{\epsilon}, \quad \boldsymbol{\epsilon} \sim \mathcal{N}(0, \mathbf{I})$$

where β_t is a **noise schedule**: $0 < \beta_1 < \beta_2 < \dots < \beta_T < 1$

Key notation: Define $\alpha_t = 1 - \beta_t$ and $\bar{\alpha}_t = \prod_{s=1}^t \alpha_s$

Closed-form: thanks to the properties of Gaussian distributions, we can analytically solve the entire sequential process (reparameterization trick):

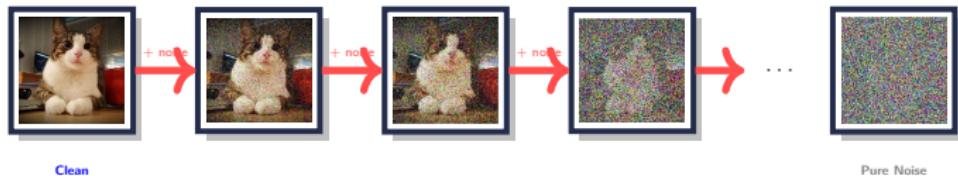
$$\mathbf{x}_t = \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \boldsymbol{\epsilon}, \quad \boldsymbol{\epsilon} \sim \mathcal{N}(0, \mathbf{I})$$

This closed form is **crucial** for efficient training!

Can jump directly to any timestep t without iterating

Diffusion Process Visualization

Forward (fixed)



Reverse (learned)



Challenge: How do we learn to reverse this process?

The Reverse Process: The Central Challenge

We have the (easy) Forward Process: We know how to add noise step-by-step: $p(\mathbf{x}_t | \mathbf{x}_{t-1})$

$$\mathbf{x}_0 \rightarrow \mathbf{x}_1 \rightarrow \cdots \rightarrow \mathbf{x}_T$$

We need the (hard) Reverse Process: To generate, we must learn to remove noise step-by-step: $p(\mathbf{x}_{t-1} | \mathbf{x}_t)$

$$\mathbf{x}_0 \leftarrow \mathbf{x}_1 \leftarrow \cdots \leftarrow \mathbf{x}_T$$

The Problem: This reverse distribution $p(\mathbf{x}_{t-1} | \mathbf{x}_t)$ is **intractable**. It's unknown and depends on the entire (unknown) data distribution.

The Key Insight: It can be shown that this difficult reverse step becomes possible *if* we can estimate one thing: the gradient of the log-probability of the noisy data distribution, $\nabla_{\mathbf{x}_t} \log p_t(\mathbf{x}_t)$.

This gradient is the key to reversing the process. It has a name...

The Score Function

Definition: For probability distribution $p(\mathbf{x})$, the **score function** is:

$$s(\mathbf{x}) = \nabla_{\mathbf{x}} \log p(\mathbf{x})$$

This is the gradient of log-probability with respect to data

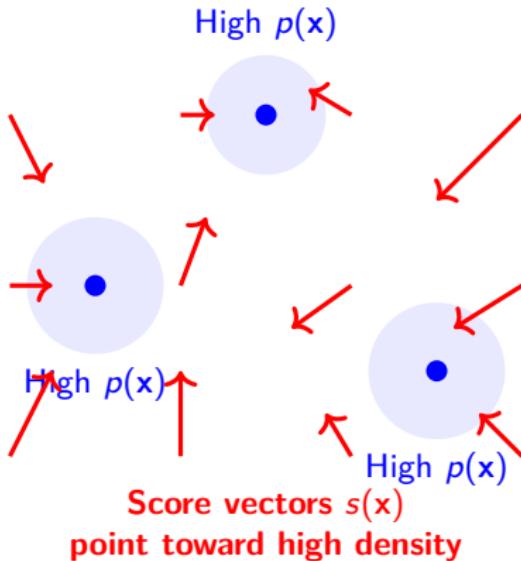
Intuition: Vector field pointing toward high-density regions

- At each point \mathbf{x} , $s(\mathbf{x})$ is a vector
- Points in direction where $\log p(\mathbf{x})$ increases most rapidly
- Following this field leads from noise to data!

For diffusion: We have score at each noise level t :

$$s_t(\mathbf{x}) = \nabla_{\mathbf{x}} \log p_t(\mathbf{x})$$

Score Function as Vector Field



Generation: Start from noise, follow the score to reach data

The Key Connection: Denoising is Score Matching

Tweedie's Formula: For noisy observation $\mathbf{x}_t = \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \boldsymbol{\epsilon}$:

$$\nabla_{\mathbf{x}_t} \log p_t(\mathbf{x}_t) = -\frac{1}{\sqrt{1 - \bar{\alpha}_t}} \mathbb{E}[\boldsymbol{\epsilon} \mid \mathbf{x}_t]$$

This means:

- The score function is proportional to expected noise
- Training to predict noise \Leftrightarrow learning the score!
- so we don't have to learn the score function directly, instead we train a neural network $\epsilon_\theta(\mathbf{x}_t, t)$ to do self-supervised noise prediction ($\boldsymbol{\epsilon}$)!
- This is called **denoising score matching**

Training objective: Train $\epsilon_\theta(\mathbf{x}_t, t)$ to predict noise:

$$\mathcal{L} = \mathbb{E}_{t, \mathbf{x}_0, \boldsymbol{\epsilon}} [\|\boldsymbol{\epsilon} - \epsilon_\theta(\mathbf{x}_t, t)\|^2]$$

where $\mathbf{x}_t = \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \boldsymbol{\epsilon}$

Why Predict Noise Instead of Images?

Alternative formulations:

- Predict clean image \mathbf{x}_0 directly
- Predict mean μ_θ directly

Why noise prediction is better:

- ➊ **Simpler objective:** Just MSE loss on noise
- ➋ **Better gradient flow:** Avoids predicting averages
- ➌ **Connection to score:** $\epsilon_\theta \approx -\sqrt{1 - \bar{\alpha}_t} \nabla_{\mathbf{x}_t} \log p_t(\mathbf{x}_t)$
- ➍ **Stationary target:** Noise is simple, stationary distribution

Once we predict noise, we can:

- Recover the score function
- Compute the denoising step to get \mathbf{x}_{t-1}

Training Procedure (Remarkably Simple)

For each training step:

- ➊ Sample data point \mathbf{x}_0 from dataset
- ➋ Sample random timestep $t \sim \text{Uniform}(1, T)$
- ➌ Sample random noise $\epsilon \sim \mathcal{N}(0, \mathbf{I})$
- ➍ Create noisy version: $\mathbf{x}_t = \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon$
- ➎ Predict noise: $\hat{\epsilon} = \epsilon_\theta(\mathbf{x}_t, t)$
- ➏ Minimize MSE: $\mathcal{L} = \|\epsilon - \hat{\epsilon}\|^2$

That's it! Just predict the noise you added

Sampling: Generating Images

To generate a new image:

- ➊ Start with pure noise: $\mathbf{x}_T \sim \mathcal{N}(0, \mathbf{I})$
- ➋ For $t = T, T - 1, \dots, 1$:
 - Predict noise: $\hat{\epsilon} = \epsilon_\theta(\mathbf{x}_t, t)$
 - Compute mean: $\mu = \frac{1}{\sqrt{\alpha_t}} \left(\mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \hat{\epsilon} \right)$
 - Sample $\mathbf{z} \sim \mathcal{N}(0, \mathbf{I})$ if $t > 1$, else $\mathbf{z} = \mathbf{0}$
 - Update: $\mathbf{x}_{t-1} = \mu + \sigma_t \mathbf{z}$
- ➌ Return \mathbf{x}_0

Intuition:

- At each step: predict and remove noise
- Add small random noise for stochasticity (except last step)
- Gradually reveal the image by following the score

From Self-Supervision to Creativity

A puzzle emerges:

- Diffusion models learn to denoise images (self-supervised)
- They're trained to predict noise as accurately as possible
- Yet they generate novel, creative images not in training data

The apparent contradiction:

- ① Perfect learning → learns ideal score function exactly
- ② Ideal score function → perfectly reverses forward process
- ③ Perfect reversal → only generates memorized training examples
- ④ But we observe: Creative, novel outputs!

The Central Puzzle

Question: How do diffusion models produce creative outputs?

Novel combinations not in training data?

The Paradox:

If model perfectly learns ideal score function on finite dataset, it can only memorize training data!

Why? For finite dataset $\mathcal{D} = \{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}\}$:

$$p_t(\mathbf{x}) = \frac{1}{N} \sum_{i=1}^N \mathcal{N}(\mathbf{x} \mid \sqrt{\bar{\alpha}_t} \mathbf{x}^{(i)}, (1 - \bar{\alpha}_t) \mathbf{I})$$

As $t \rightarrow 0$, posterior concentrates on nearest training image

Perfect training = only memorization

Again, where does creativity come from?

Creativity as Structured Failure

Theoretical Insight: Creativity arises because model *fails* to learn ideal score

Crucially: This failure is **structured by inductive biases!**

For CNN-based U-Net, two biases are key:

① **Locality:** Finite receptive fields

- Score at pixel (i, j) depends only on local neighborhood
- Cannot coordinate globally instantaneously
- **Connects to self-supervision:** Each patch makes independent denoising decisions

② **Equivariance:** Weight sharing (Lecture 4!)

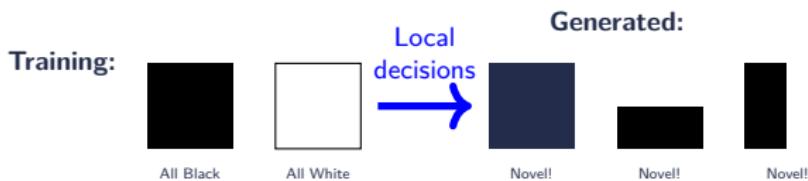
- CNNs treat different locations similarly
- Translation invariance
- **Connects to self-supervision:** Denoising strategy learned on one patch applies everywhere

These architectural constraints prevent implementing an “Ideal Score Machine”

Result: The model denoises locally, composing globally novel mosaics.

The Simplest Example: Black and White Images

Training set: Only 2 images (all black, all white)



Exponentially many novel samples! (approximately 2^{N^2} for $N \times N$ image)

How? Each pixel independently decides its color based on local neighborhood

Local consistency: majority color in patch = center pixel

The Mechanism: Patch Mosaics

What happens instead of memorization:

① Local Bayesian Inference:

- Each pixel estimates local score using only nearby info
- “Which training patch do I most resemble?”

② Mixing and Matching:

- Model doesn't memorize whole images
- Composes patches from different training images

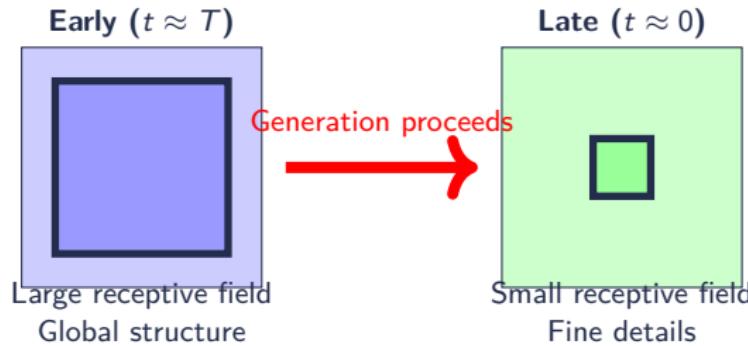
③ Locally Consistent, Globally Novel:

- Every small patch looks realistic (matches training)
- Overall combination is new (never seen)
- Combinatorial creativity!

Coarse-to-Fine Generation

Important empirical observation:

Effective receptive field shrinks during reverse process



Strategy:

- **Early:** Large patches set global structure (object type, layout)
- **Late:** Small patches add fine details (textures, edges)

Explaining Spatial Inconsistencies

Famous diffusion “errors”:

- Hands with wrong number of fingers
- Clothing with incorrect number of arms
- Bifurcated shoes or multiple legs on pants

Mechanistic explanation: Excessive locality at late times ($t < 0.3$)

- Receptive field < 5 pixels
- Different parts of image cannot coordinate
- Each region independently decides “this should be a finger”
- Result: Too many fingers!

This is not a bug—it's a fundamental consequence
of the local score approximation that enables creativity

It's a trade-off!

Connection to Lecture 6: U-Net Returns!

Recall from Lecture 6: U-Net for image segmentation

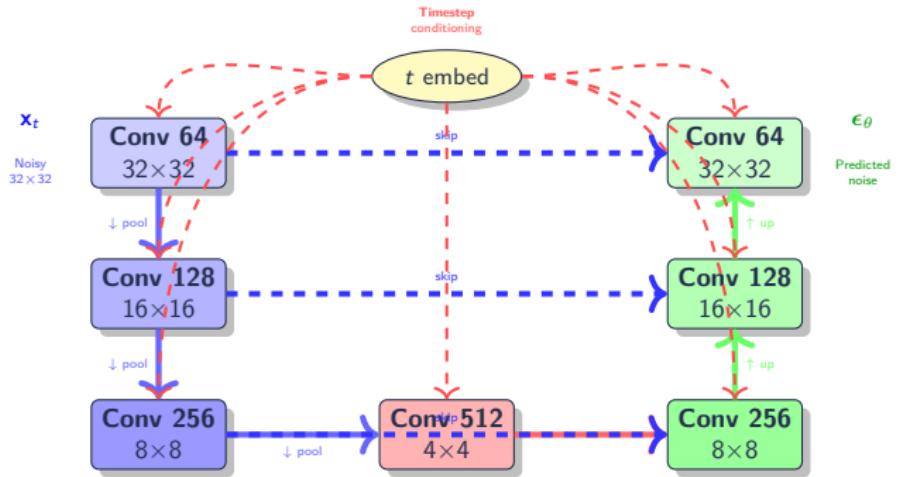
Perfect for diffusion because:

- **Spatial structure:** Preserves image layout
- **Multi-scale:** Handles coarse and fine details
- **Skip connections:** Essential for preserving details during denoising
- **Image-to-image:** Noisy image → noise prediction

Typical architecture: $\epsilon_\theta(\mathbf{x}_t, t)$

- **Input:** Noisy image \mathbf{x}_t + timestep t
- **Output:** Predicted noise $\hat{\epsilon}$
- **Timestep embedding:** Sinusoidal encoding (like Transformers!)

U-Net for Diffusion



Skip connections are crucial: preserve high-frequency details lost in downsampling

The Challenge: Pixel-Space is Expensive

DDPM works, but:

- 512×512 image = 786,432 pixels
- Need 50-1000 denoising steps
- Running U-Net 1000 times on 512×512 is prohibitive!

Key Observation: Most image information is redundant!

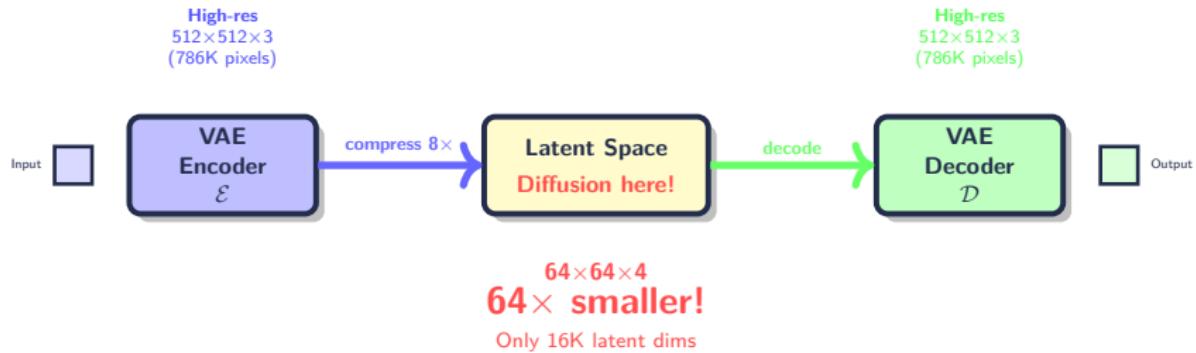
Nearby pixels are highly correlated

Solution: Latent Diffusion

Run diffusion in compressed latent space

Latent Diffusion Models (LDM)

Idea: Use pre-trained VAE to compress images



Process:

- ① Encode image to latent: $\mathbf{z} = \mathcal{E}(\mathbf{x})$
- ② Run diffusion on \mathbf{z} (much smaller!)
- ③ Decode back: $\hat{\mathbf{x}} = \mathcal{D}(\mathbf{z})$

Benefits of Latent Diffusion

Why this is game-changing:

- **Speed:** 64×64 latent vs 512×512 pixels
 - $\frac{512^2}{64^2} = 64 \times$ fewer pixels per step!
 - Same quality, dramatically faster
- **Memory:** Can train on consumer GPUs
 - 512×512 diffusion: needs A100 (80GB)
 - 64×64 latent: works on RTX 3090 (24GB)
- **Quality:** Still generate high-res images
 - VAE decoder upsamples from latent
 - Preserves details surprisingly well

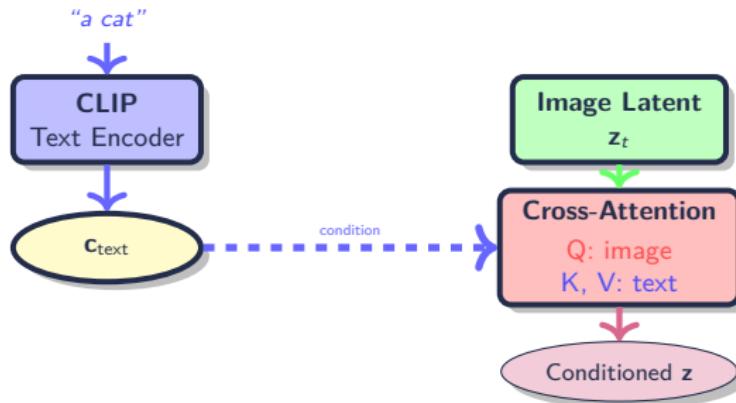
This is what Stable Diffusion uses!

Text Conditioning via Cross-Attention

Challenge: Generate specific content, not random images

Need to condition on text prompts!

Solution: Cross-attention between image and text



Mechanism: Image patches “query” text to find relevant semantic info

Cross-Attention Mechanism

Recall from Lecture 8: Attention mechanism

$$\text{Attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax}\left(\frac{\mathbf{Q}\mathbf{K}^\top}{\sqrt{d}}\right)\mathbf{V}$$

For cross-attention in diffusion:

- $\mathbf{Q} = \mathbf{W}_Q \cdot \mathbf{z}_t$ (query from noisy image)
- $\mathbf{K} = \mathbf{W}_K \cdot \mathbf{c}_{\text{text}}$ (key from CLIP text)
- $\mathbf{V} = \mathbf{W}_V \cdot \mathbf{c}_{\text{text}}$ (value from CLIP text)

Intuition:

When generating cat's ear, image latent "attends to" "cat" in text embedding

Each image region focuses on relevant text concepts

Classifier-Free Guidance (CFG)

Problem: Text conditioning alone may be too weak

Solution: Amplify the conditioning effect!

Training: Randomly drop text 10% of time

- Model learns both $\epsilon_\theta(\mathbf{x}_t, t, \mathbf{c})$ and $\epsilon_\theta(\mathbf{x}_t, t, \emptyset)$

Sampling: Use guided prediction

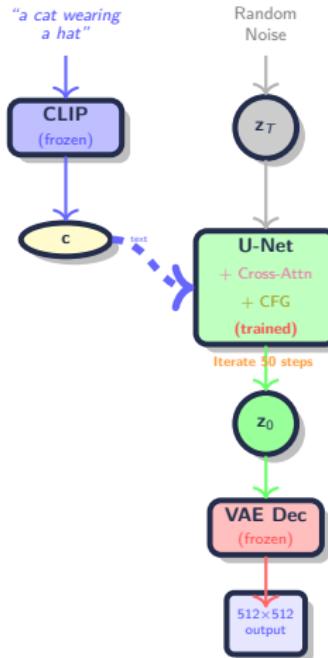
$$\tilde{\epsilon} = \epsilon_\theta(\mathbf{x}_t, t, \emptyset) + s \cdot [\epsilon_\theta(\mathbf{x}_t, t, \mathbf{c}) - \epsilon_\theta(\mathbf{x}_t, t, \emptyset)]$$

where $s > 1$ is guidance scale (typically 7.5)

Intuition: Move away from unconditional, toward conditional

Higher $s \rightarrow$ stronger conditioning but less diversity

Complete Stable Diffusion Pipeline



Components: CLIP (frozen) + U-Net (trained) + VAE (frozen)

Faster Sampling

Problem: DDPM needs 1000 steps, too slow!

Method	Steps	Description
DDIM	50-100	Deterministic, skip steps
DPM-Solver	20-50	ODE solver for diffusion
Flow Matching	10-20	Continuous flows
Consistency	1-4	Direct mapping

DDIM: Deterministic sampling with fewer steps

Make sampling deterministic, skip timesteps: 50 steps achieves similar quality

Diffusion Transformers (DiT)

Recent trend: Replace U-Net with Transformer blocks

DiT architecture:

- ① **Patchify:** Split latent into patches (like ViT!)
- ② **Add embeddings:** Position + timestep
- ③ **Transformer blocks:** Multi-head attention + FFN
- ④ **Unpatchify:** Reshape to latent space

Advantages over U-Net:

- Scalability: Easy to scale up
- Long-range dependencies: Global self-attention
- Unified architecture: Same as ViT, BERT, GPT

Results: DiT matches or exceeds U-Net quality with better scaling!

Evaluation Metrics

How do we measure quality?

Metric	What it measures	Range
FID	Distribution similarity	Lower better
CLIP Score	Text-image alignment	Higher better
Human Eval	Preference ratings	Subjective

FID (Fréchet Inception Distance):

- Compare feature distributions of real vs generated
- Extract features using Inception-v3
- Fit Gaussian, compute Fréchet distance
- Lower FID = closer to real data

What We've Learned: The Journey

1. The Averaging Problem:

- Predicting multiple correlated values → blurring
- Solution: Predict one at a time (auto-regressive)

2. The Insight: Break correlations with noise

- Add independent Gaussian noise to all pixels
- Can predict all simultaneously without averaging!

3. Mathematical Foundation:

- Score function $s(\mathbf{x}) = \nabla_{\mathbf{x}} \log p(\mathbf{x})$ guides generation
- Denoising = Score matching (Tweedie's formula)
- Training: Just predict noise with MSE loss

What We've Learned: The Practice

4. Architecture: U-Net from Lecture 6

- Skip connections preserve details
- Multi-scale processing
- Timestep conditioning

5. Creativity Paradox:

- Perfect learning → memorization
- Creativity from structured failure
- CNN locality + equivariance → patch mosaics
- Coarse-to-fine: large patches (early) to small (late)

6. Stable Diffusion: Making it practical

- Latent diffusion ($64\times$ faster)
- Cross-attention for text conditioning
- Classifier-free guidance (amplify conditioning)

Key Takeaways

① Diffusion revolutionized image generation

- Stable training, high quality, excellent diversity
- Solved problems that plagued GANs

② Self-supervised learning is key

- No labels needed, just images
- Denoising provides natural training signal

③ Architecture matters

- U-Net perfect for image-to-image tasks
- Inductive biases shape creativity
- Future: Transformer-based (DiT)

④ Efficiency through clever design

- Latent space diffusion ($64\times$ speedup)
- Cross-attention for conditioning
- Classifier-free guidance for control

The Generative AI Revolution

From 2020 to 2025:

- 2020: DDPM introduces stable training
- 2021: DALL-E shows text-to-image, CLIP enables grounding
- 2022: Stable Diffusion democratizes (open source!)
- 2023: Midjourney reaches photorealism, video begins
- 2024: Sora generates minute-long videos
- 2025: Multimodal unified models

Applications:

- Art, design, advertising
- Scientific research (proteins, drugs)
- Text-to-video (Sora, Runway)
- Image editing, super-resolution

We're still in the early days!