


Lecture 12.3: GenAI: Variational Autoencoders

Heman Shakeri

Recall: The Autoencoder from Module 6

We've seen autoencoders as unsupervised learning tools:



Goal: Minimize reconstruction error $\|\mathbf{x} - \hat{\mathbf{x}}\|^2$

Latent code z : Compressed representation of input


The Problem: Can we generate new samples?

Try sampling random z and decoding...

Recall the Latent Space Geometry


What we hope for:

- Continuous latent space
- Smooth interpolation
- Every point decodes to something meaningful



What we actually get:

- Scattered, disconnected regions
- Random points \rightarrow garbage
- No principled way to sample



We need the latent space to be a *continuous probability distribution!*

VAE

The Paper That Started It All

“Auto-Encoding Variational Bayes” (2013)

Diederik P. Kingma and Max Welling

<https://arxiv.org/abs/1312.6114>

Idea: Instead of learning a single point \mathbf{z} for each input, learn a *distribution* over \mathbf{z} .


Specifically: Learn parameters μ and σ of a Gaussian

$$q(\mathbf{z}|\mathbf{x}) = \mathcal{N}(\mathbf{z}; \boldsymbol{\mu}(\mathbf{x}), \text{diag}(\boldsymbol{\sigma}^2(\mathbf{x})))$$

Why this is genius: We can now sample from the latent space to generate!

VAE Architecture

Key difference from standard AE: Encoder outputs distribution parameters, not a point!



The Solution: A Lower Bound

Strategy: Use a Tractable Proxy

Since we can't maximize $\log p(\mathbf{x})$, we find a new, tractable function (the ELBO) that is a lower bound. By

maximizing this proxy, we push up the true likelihood.

The Jensen's Inequality Trick

We use it because \log is a **concave** function, which means:

$\log(\mathbb{E}[Y]) \geq \mathbb{E}[\log(Y)]$. This one rule lets us create the bound.

Terse Derivation

- Start with $\log p(\mathbf{x})$ and introduce $q(\mathbf{z}|\mathbf{x})$:

$$= \log \int q(\mathbf{z}|\mathbf{x}) \frac{p(\mathbf{x}, \mathbf{z})}{q(\mathbf{z}|\mathbf{x})} d\mathbf{z}$$

- Rewrite as an expectation:

$$= \log \left(\mathbb{E}_q \left[\frac{p(\mathbf{x}, \mathbf{z})}{q(\mathbf{z}|\mathbf{x})} \right] \right)$$

- Apply Jensen's (move \log inside):

$$\geq \mathbb{E}_q \left[\log \left(\frac{p(\mathbf{x}, \mathbf{z})}{q(\mathbf{z}|\mathbf{x})} \right) \right]$$

- Rearrange... and we get the ELBO!

The Beautiful Math: ELBO

The Challenge: We want to maximize $\log p(\mathbf{x})$ (likelihood of our data), but it's intractable!

$$\log p(\mathbf{x}) = \log \int p(\mathbf{x}|\mathbf{z})p(\mathbf{z})d\mathbf{z}$$

Solution: Introduce approximate posterior $q(\mathbf{z}|\mathbf{x})$ and derive a lower bound
Using Jensen's inequality, we get the **Evidence Lower BOund (ELBO)**:

$$\log p(\mathbf{x}) \geq \mathbb{E}_{q(\mathbf{z}|\mathbf{x})}[\log p(\mathbf{x}|\mathbf{z})] - D_{KL}(q(\mathbf{z}|\mathbf{x})\|p(\mathbf{z}))$$

This is what we maximize!

Two components:

- ① **Reconstruction term:** How well can we reconstruct \mathbf{x} from \mathbf{z} ?
- ② **Regularization term:** How close is our posterior to the prior?

Understanding the ELBO Intuitively

$$\mathcal{L} = \underbrace{\mathbb{E}_{q(z|x)}[\log p(x|z)]}_{\text{Reconstruction term}} - \underbrace{D_{KL}(q(z|x) \| p(z))}_{\text{keep the latent space tidy!}}$$

Reconstruction term: Standard autoencoder objective

- “Can I decode z back to x ? ”
- Encourages preserving information

KL divergence term: The magic ingredient!

- “Is $q(z|x)$ close to standard normal $\mathcal{N}(0, I)$? ”
- Forces latent codes to be well-behaved
- Prevents encoder from cheating by using arbitrary regions
- Creates continuous, smooth latent space

The trade-off: Balance reconstruction quality vs. latent space structure

The KL Divergence: Forcing Structure

For Gaussians, the KL divergence has a closed form!

$$D_{KL}(q(\mathbf{z}|\mathbf{x})\|p(\mathbf{z})) = \frac{1}{2} \sum_{j=1}^J (\mu_j^2 + \sigma_j^2 - \log \sigma_j^2 - 1)$$

where J is the latent dimension.

What does this do?

- Pulls μ toward zero
- Pulls σ toward one
- Prevents collapse to deterministic encoding

Result: All latent codes live in a similar region around $\mathcal{N}(0, I)$

This means random samples from $\mathcal{N}(0, I)$ will decode to valid outputs!


The Reparameterization Trick: Visual

Problem: We need to sample $\mathbf{z} \sim q(\mathbf{z}|\mathbf{x})$, but backpropagation cannot flow through a random sampling operation.

$$\mathbf{z} \sim q(\mathbf{z}|\mathbf{x}) = \mathcal{N}(\mu, \sigma^2)$$

Sampling is *not* differentiable! \rightarrow Reparameterize the sampling:
Instead of sampling \mathbf{z} directly, write:

$$\mathbf{z} = \mu + \sigma \odot \epsilon, \quad \text{where } \epsilon \sim \mathcal{N}(0, I)$$



[Kingma, 2013; Bengio, 2013; Rezende et al 2014]

β -VAE: Forcing Disentanglement

- A standard VAE ($\beta = 1$) has a problem.
- The reconstruction term often “wins” the trade-off, forcing the model to be a perfect autoencoder.
- The result: it ignores the KL term, creating a messy, **entangled** latent space just to pass information.

What is Disentanglement?

We want each latent dimension z_i to control *one single, independent* factor of the data.

Example:

- For faces, z_1 might control “smile,” z_2 controls “head rotation,” and z_3 controls “skin tone.”
- A simple VAE fails at this; z_1 might control both smile *and* rotation (entangled).

The β -VAE Solution

We add a simple hyperparameter β to the KL term:

$$\mathcal{L} = \mathbb{E}_q[\log p(\mathbf{x}|\mathbf{z})] - \beta \cdot D_{KL}(q||p)$$

Intuition: When $\beta > 1$, we put *more pressure* on the model to be structured and **forced** to find the most efficient encoding: the true, underlying, independent factors. This is disentanglement.

PyTorch Implementation

```
1 class VAE(nn.Module):
2     def __init__(self, input_dim, hidden_dim, latent_dim):
3         super().__init__()
4         # Encoder
5         self.encoder = nn.Sequential(
6             nn.Linear(input_dim, hidden_dim),
7             nn.ReLU(),
8             nn.Linear(hidden_dim, hidden_dim),
9             nn.ReLU()
10        )
11        self.fc_mu = nn.Linear(hidden_dim, latent_dim)
12        self.fc_logvar = nn.Linear(hidden_dim, latent_dim)
13
14        # Decoder
15        self.decoder = nn.Sequential(
16            nn.Linear(latent_dim, hidden_dim),
17            nn.ReLU(),
18            nn.Linear(hidden_dim, hidden_dim),
19            nn.ReLU(),
20            nn.Linear(hidden_dim, input_dim),
21            nn.Sigmoid()
22        )
```

VAE Forward Pass and Loss

```
1 def encode(self, x):
2     h = self.encoder(x)
3     mu = self.fc_mu(h)
4     logvar = self.fc_logvar(h)
5     return mu, logvar
6
7 def reparameterize(self, mu, logvar):
8     std = torch.exp(0.5 * logvar) # sigma = exp(0.5 * log(sigma^2))
9     eps = torch.randn_like(std) # Sample epsilon ~ N(0,1)
10    return mu + eps * std # z = mu + sigma * epsilon
11
12 def forward(self, x):
13     mu, logvar = self.encode(x)
14     z = self.reparameterize(mu, logvar)
15     recon_x = self.decoder(z)
16     return recon_x, mu, logvar
17
18 def vae_loss(recon_x, x, mu, logvar):
19     # Reconstruction loss (binary cross-entropy)
20     recon_loss = F.binary_cross_entropy(recon_x, x, reduction='sum')
21     # KL divergence
22     kl_loss = -0.5 * torch.sum(1 + logvar - mu.pow(2) - logvar.exp())
23     return recon_loss + kl_loss
```

Generation: Sampling from the Prior

Training: Encode data → sample latent → decode

Generation: Skip the encoder!

- ① Sample $\mathbf{z} \sim \mathcal{N}(0, I)$ directly from prior
- ② Pass through decoder: $\hat{\mathbf{x}} = \text{Decoder}(\mathbf{z})$
- ③ Get a new sample!



Why this works: KL term forced all encodings near $\mathcal{N}(0, I)$


Random samples from this region decode to valid data!

Interpolation: The Power of Continuous Space

Because latent space is continuous, we can interpolate!

Procedure:

- ① Encode two images: $\mathbf{z}_1 = \text{Encode}(\mathbf{x}_1)$, $\mathbf{z}_2 = \text{Encode}(\mathbf{x}_2)$
- ② Interpolate: $\mathbf{z}_t = (1 - t)\mathbf{z}_1 + t\mathbf{z}_2$ for $t \in [0, 1]$
- ③ Decode: $\mathbf{x}_t = \text{Decode}(\mathbf{z}_t)$



This doesn't work with standard autoencoders because their latent space has holes!

Connection to Diffusion Models

Remember Latent Diffusion from the previous lecture?

Stable Diffusion uses a VAE!

- Pre-trained VAE encoder compresses images $8\times$
- Diffusion happens in latent space (64×64 instead of 512×512)
- VAE decoder upsamples final result

Why VAE specifically?

- Continuous latent space \rightarrow perfect for diffusion
- Learned compression preserves important features
- Separate the compression problem from generation

This is a brilliant example of composing different generative models!

Limitations of VAEs

Blurriness problem:

- VAE samples tend to be blurrier than GANs
- Why? Reconstruction loss averages over possibilities
- The prior assumption (Gaussian) may be too restrictive

Posterior collapse:

- Sometimes decoder ignores latent code
- KL term goes to zero, no information in \mathbf{z}
- Solutions: Annealing β , architectural changes

Difficulty with complex distributions:

- Gaussian assumption may be too simple
- Real data distributions are multimodal, complex
- Led to variants: VQ-VAE, Normalizing Flows, etc.

VAE Variants and Extensions

Conditional VAE (CVAE):

- Condition on class labels or attributes
- Control what to generate

VQ-VAE (Vector Quantized VAE):

- Discrete latent space instead of continuous
- Used in DALL-E, better for high-quality images
- Learned codebook of latent vectors

Hierarchical VAE:

- Multiple levels of latent variables
- Captures structure at different scales

Importance Weighted AE (IWAE):

- Tighter bound on log-likelihood
- Better density estimation

Why VAEs Matter

Theoretical elegance:

- Principled probabilistic framework
- Interpretable objective (ELBO)
- Connections to information theory, Bayesian inference

Practical applications:

- Image generation and compression
- Anomaly detection (reconstruction error)
- Representation learning for downstream tasks
- Semi-supervised learning
- Data imputation and denoising

Foundation for modern generative models:

- Ideas appear in diffusion models
- VQ-VAE powers DALL-E
- Latent space manipulation techniques

Key Takeaways

1. The Core Idea:

- Encode data as distributions, not points
- Use reparameterization trick for backprop

2. The ELBO:

- Reconstruction: preserve information
- KL divergence: regularize latent space

3. The Power:

- Principled generation by sampling from prior
- Continuous latent space enables interpolation
- Probabilistic framework with theoretical guarantees

4. The Legacy:

- Foundation for modern generative AI
- Still used in Stable Diffusion today
- Inspired countless variants and improvements